


Today's AI safety stack was built for models that answer, not agents that act.



This protects conversations. It does not fully govern actions.


AUTONOMY SHIFT

As agents gain autonomy, the safety problem becomes authority control.

Agents are moving from answering questions to operating across enterprise systems.




The more authority an agent has, the less prompt-only safety scales.



FAILURE MODE

The critical failure is when untrusted input causes privileged action.




The risk is that language becomes **authority**.

TO-BE SAFETY ARCHITECTURE

Future guardrails must control the full path from intent to egress.

A single output filter cannot govern multi-step agent behavior.



Not just:

Is this output safe?



But:

Was this action authorized by the original intent?

ScopeTrace makes agent guardrails traceable, enforceable, and replayable.

